

Sprint Retrospective - 1

Date: 02/03/2023

Authors: Hongbin, Michael, Luis

Introduction:

This retrospective will serve as documentation for this latest sprint, as a way to catalogue the process.

Good Things:

- Most error found quickly and was easy to fix
- Everyone understood what their designated task was during any given incident
- Communication between members was great
- All planned user stories were completed

Opportunities for Improvement:

- Get better estimates for how long user stories will take to complete
- Faster image rollback upon bug discovery

Next Sprint Planning:

Two epics will be implemented next sprint, those being separating the monitoring tools and application pods and hosting them on different nodes as well as configuring Prometheus to send alerts regarding cluster conditions that may prevent SLOs from being met. The separation of the application pods from the monitoring tools will ensure cascading failures in one will not affect the other. Additional Prometheus alerts will ensure issues such as high memory usage or low file space will be detected early and not lead to node failures. The velocity for each of those epics to be fully implemented should be two days.

Wrap-Up:

Did well on the actual debugging and fixing of said bugs, but we could have a better reaction time for rolling back our images once a problem is detected. While we did have some spare time after the two planned stories were implemented, we don't believe that a third full story would have been doable. Instead, we started preparatory work for the next sprint by adjusting our manifest files to include cluster metrics in the Grafana dashboard.